

# Neural Networks

## Introduction to Deep Learning

# Agenda

- Introduction to neural networks
  - Neural Networks
  - Neurons
- Activation functions
  - Sigmoid, Tanh, ReLU
- Feed Forward neural network
  - Layer Details
- Training a neural network

# Agenda

- Error and Loss function
- Optimization
- Gradient descent
  - Gradient
  - Gradient Descent Variations
  - Backpropagation
- Summary





# Neural Networks

*Artificial Neural Networks are computing systems inspired from biological neuron*

# Neuron

Artificial neuron is inspired by biological neuron

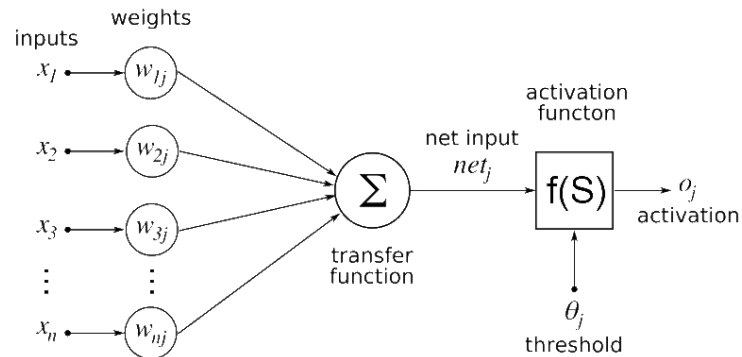


# Activation function


In artificial neural networks, helps in defining the output of a node when a input is given.

Different types of activations helps in different tasks. Some examples are -


- Sigmoid
- Tanh
- ReLU
- Binary Step Function




# Types of activation function



Sigmoid



ReLU



Tanh

Activation function performs certain mathematical operations on its input, which is a number



Question:

What is the range of Sigmoid, Tanh and ReLU?

Answer:

Sigmoid (0,1)

Tanh (-1, 1)


ReLU (0, max)






Activation function takes in the output signal from the previous cell and converts it into some form that can be taken as input to the next cell.

# Feed forward neural network



This is a 2-layer neural network. One is the hidden layer (having 3 neurons) and the other is output layer (having 2 neurons).

# Layer details



## Output layer

- Represents the output of the neural network

## Hidden layer(s)

- Represents the intermediary nodes.
- It takes in a set of weighted input and produces output through an activation function

## Input layer


- Represents dimensions of the input vector (one node for each dimension)

# Training a neural network

- Decide the structure of network
- Create a neural network
- Choose different hyper-parameters
- Calculate loss
- Reduce loss
- Repeat last three steps

# Error and Loss function

- In general, error/loss for a neural network is difference between actual value and predicted value.
- The goal is to minimize the error/loss.
- Loss Function is a function that is used to calculate the error.
- You can choose loss function based on the problem you have at hand.
- Loss functions are different for classification and regression




# Optimization

In optimization, the main aim is to find weights that reduce loss


# Gradient

- Gradient is calculated by optimization function
- Gradient is the change in loss with change in weights.
- The weights are modified according to the calculated gradient.
- Same process keep on repeating until the minima is reached



# Gradient descent

Gradient descent is a method that defines a cost function of parameters and uses a systematic approach to optimize the values of parameters to get minimum cost function.







Question:

Do we get multiple local optimum solutions if we solve using gradient descent.

Answer:

False. We get only one local optimum solution after using gradient descent

# Gradient descent variations

Gradient descent has 3 variations, these differ in using data to calculate the gradient of the objective function

1. Batch gradient descent
  - Updates the parameter by calculating gradients of whole dataset
2. Stochastic gradient descent
  - Updates the parameters by calculating gradients for each training example
3. Mini -batch gradient descent
  - Updates the parameters by calculating gradients for every mini batch of "n" training example
  - Combination of batch and stochastic gradient descent

# Backpropagation

- Backpropagation is used while training the feedforward networks
- It helps in efficiently calculating the gradient of the loss function w.r.t weights
- It helps in minimizing loss by updating the weights

# Learning rate and Momentum

- The learning rate is a hyperparameter which determines to what extent newly acquired weights overrides old weights. In general it lies between 0 and 1.
- You can try different learning rates for a neural networks to improve results.
- Momentum is used to decide the weight on nodes from previous iterations. It helps in improving training speed and also in avoiding local minimas.

# Summary of the Neural Network Process

- A neural network is made of neurons
- These neurons are connected to each other
- Every neuron has an activation function that defines its output
- Then we train our neural network to learn the parameter values i.e. weights and biases
- This process consists of forwardprop and backprop
- After forward prop we calculate the loss using a loss function and propagate the information backwards, that's backprop
- This process is repeated layer by layer, until all the neurons receive a loss signal which describes their contribution to the total loss

# TensorFlow

## Introduction

# Agenda

- What is tensorflow?
- Why we are using tensorflow?
- TensorFlow 2.x
- TensorFlow 2.x - features and changes
- Changes with respect to TensorFlow 1.x
- Keras and it's advantages
- Keras vs tf.keras
- Tutorials and guides
- Notes

# TensorFlow

- TensorFlow is a machine learning library by Google
- It is open sourced
- It is mainly used for implementing neural networks
- It is an end-to-end platform, which means you can use it for building your models from scratch to deploying them into a production environment



# Why TensorFlow for this course?

- Most used library for deep learning
- Easy transition from research to production
- Extensive industry support
- Easy deployment around servers, mobile devices, web platforms etc
- You can easily sort out the issues using GitHub, StackOverflow etc
- Used by world's top AI companies
- Consistently updated with cutting edge changes

# Some technical reasons

- Easy and readable syntax
- Being a low-level library it provide more flexibility to developers to implement their own functionalities and services
- Provides high level API for implementing advanced neural net architectures
- Distributed training

# OS Support

TensorFlow is supported on following 64-bit systems:

- Ubuntu 16.04 or later
- macOS 10.12.6 (Sierra) or later (no GPU support)
- Windows 7 or later
- Raspbian 9.0 or later

# Language support

- Python
- C++
- JavaScript
- Java
- Go
- Swift

# What is a Tensor?

- A tensor is an n-dimensional array which can be a scalar, vector or matrix etc
- Tensorflow uses data flow graphs for parallel computing
- Data flow graphs provide parallelism, distributed execution, faster compilation and portability
- Tensorflow can be used for developing cool projects like Image classification, speech recognition, object detection, transfer learning etc.

# Tensors

- A Tensor is a multi-dimensional array.
- Similar to NumPy `ndarray` objects, `tf.Tensor` objects have a data type and a shape. Additionally, `tf.Tensor`s can reside in accelerator memory (like a GPU).
- TensorFlow offers a rich library of operations (`tf.add`, `tf.matmul`, `tf.linalg.inv` etc.) that consume and produce `tf.Tensor`s.
- These operations automatically convert native Python types, for example:

```
print(tf.add(1, 2))
print(tf.add([1, 2], [3, 4]))
print(tf.square(5))
print(tf.reduce_sum([1, 2, 3]))

# Operator overloading is also supported
print(tf.square(2) + tf.square(3))
```

# Tensor vs Numpy

The most obvious differences between NumPy arrays and `tf.Tensors` are:

1. Tensors can be backed by accelerator memory (like GPU, TPU).
2. Tensors are immutable.

## NumPy Compatibility

Converting between a TensorFlow `tf.Tensor` and a NumPy `ndarray` is easy:

- TensorFlow operations automatically convert NumPy `ndarrays` to Tensors.
- NumPy operations automatically convert Tensors to NumPy `ndarrays`.

Tensors are explicitly converted to NumPy `ndarrays` using their `.numpy()` method.

# GPU Acceleration

- Many TensorFlow operations are accelerated using the GPU for computation.
- Without any annotations, TensorFlow automatically decides whether to use the GPU or CPU for an operation—copying the tensor between CPU and GPU memory, if necessary.
- Tensors produced by an operation are typically backed by the memory of the device on which the operation executed, for example:

```
x = tf.random.uniform([3, 3])

print("Is there a GPU available: "),
print(tf.config.experimental.list_physical_devices("GPU"))

print("Is the Tensor on GPU #0: "),
print(x.device.endswith('GPU:0'))
```

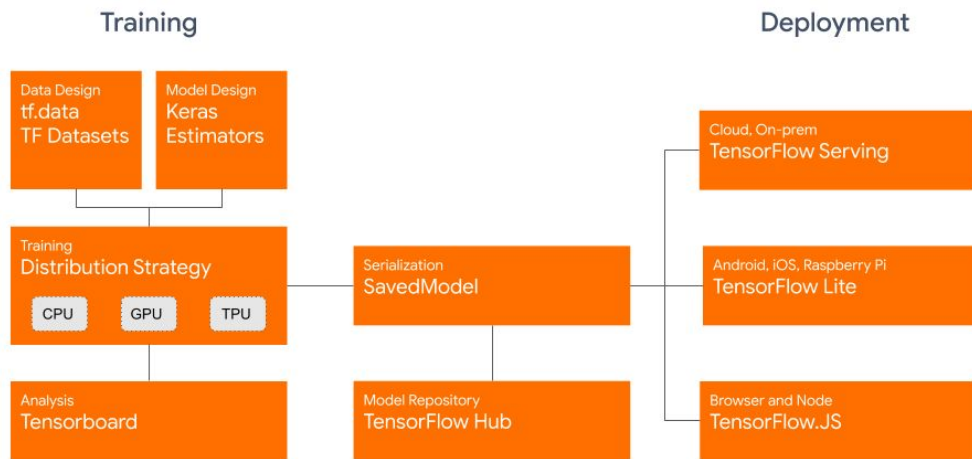
```
Is there a GPU available:
[PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
Is the Tensor on GPU #0:
True
```





# Coding with Tensorflow

- TensorFlow 2.0 makes development of DL applications much easier.
- With tight integration of Keras into TensorFlow, eager execution by default, and Pythonic function execution, TensorFlow 2.0 makes the experience of developing applications as familiar as possible for Python developers.



[Source](#)

# Basics with TensorFlow

1. Import tensorflow library

```
import tensorflow as tf
```

2. Load/Read the data

```
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
```

3. Build your model

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10)
])
```

4. Compile the model with optimizer, loss function and error metric for back propagation

```
model.compile(optimizer='adam',
              loss=loss_fn,
              metrics=['accuracy'])
```

5. Fit your model on your train data

```
model.fit(x_train, y_train, epochs=5)
```

# New version - TensorFlow 2.x

- Much easier development
- Tight integration with Keras
- More familiar for Python developers
- Deploy anywhere across servers, mobile and edge devices, browser and Node.js with TensorFlow Extended, TensorFlow Lite and TensorFlow .JS
- Multi GPU support

# What's new in TF2?

- `tf.function()`: Creates a callable TensorFlow graph from a Python function
- `tf.GradientTape()`: Records operations for automatic differentiation
- `tf.data()`: helps in building complex input pipelines from simple, and reusable pieces

# Changes with respect to TensorFlow 1.x

- Cleaning up of APIs: many APIs are either gone or moved
- Eager execution (default): no need manually compile the abstract syntax tree
- No more globals: no need to track variables
- Introduction of function and elimination of session

# Model building: from simple to flexible

- |  |        |                                  |
|--|--------|----------------------------------|
| 1. Sequential API + built in layers                                      | -----> | New users                        |
| 2. Functional API + built in layers                                      | -----> | Engineers with standard use case |
| 3. Functional API + custom layers<br>+ custom metrics<br>+ custom losses | -----> | Engineers requiring control      |
| 4. Subclassing: write everything<br>yourself from<br>scratch             | -----> | Researchers                      |

# Keras



# Keras

- Keras is a high-level neural network API
- Written and implemented in Python
- Can run on top of TensorFlow
- It was designed keeping fast experimentation in mind.

# Keras advantages

- User-friendly
  - Simple and user friendly interface. Actionable feedbacks are also provided.
- Modular and composable
  - Modules are there for every step, you can combine them to build solutions.
- Easy to extend
  - Gives freedom to add custom blocks to build on new ideas.
  - Custom layers, metrics, loss functions etc. can be defined easily.

# Keras vs tf.keras

- In TF2 instead of writing "import keras" you will write "from tensorflow import keras"
- In colab, if you see the message "Using TensorFlow Backend", you are not using tensorflow 2.x implementation of keras
- tf.keras is the TensorFlow's implementation of keras so it supports all the latest changes in the TensorFlow version
- tf.keras is also better in support and maintenance

# Latest tutorials and guides

- <https://www.tensorflow.org/tutorials>
- <https://www.tensorflow.org/guide>

# Notes

- TensorFlow 2.0 announcement video:  
<https://www.youtube.com/watch?v=EgWsPO8DVXk>
- Guide to convert your code from TF 1 to TF 2:  
<https://www.tensorflow.org/guide/migrate>
- Detailed introduction video:  
<https://www.youtube.com/watch?v=5ECD8J3dvDQ>

# Thank you!

Happy Learning :)